

# DBCC CHECKDB - Use and Abuse

## [CA Mainframe 2.0 Database](#)

Monitor & Improve DB Performance Across Subsystems, Objects & SQL.  
[www.ca.com/mainframe2](http://www.ca.com/mainframe2)

## [Business Analytics Demo](#)

See business analytics in action. See the Free Demo Now!  
[www.Cognos.com/powerplay](http://www.Cognos.com/powerplay)

## [Sql Server](#)

Ready to learn SQL Server? Try our free videos before you buy!  
[cbt nuggets.com](http://cbt nuggets.com)

## [Oracle Database Forum](#)

Discuss Oracle Database Issues, Errors, Commands, and more.  
[Oracle.ittoolbox.com](http://Oracle.ittoolbox.com)

Ads by Google

I am in the process of recruiting a senior DBA for my current client, and I am surprised at the lack of knowledge about database corruption and how to deal with it!

It has prompted me to write this article.

The question I ask in an interview is:

*"You get an error in the logs indicating that you have a torn page or checksum error. What do you do?"*

Half the candidates look confused. Of the rest, roughly 90% tell me to run DBCC CHECKDB with one of the repair options! When I challenge them to explain how CHECKDB performs the repair, none are able to answer, unaware of the damage they might be doing.

So here is how to use DBCC CHECKDB, and what to do when you have a torn or corrupt page.

## So How Do I Use It?

The primary purpose is to check for consistency errors, and should ideally be run every day.

The basic syntax is:

```
DBCC CHECKDB ('DB Name') WITH NO_INFOMSGS
```

NO\_INFOMSGS prevents an excessive number of informational messages from being generated. There are several other options, but this is the syntax you should aim to use as it performs all integrity checks.

This may take a long time on large databases and you may want to specify the PHYSICAL\_ONLY option. This checks physical on-disk structures, but omits the internal logical checks. The syntax is:

```
DBCC CHECKDB ('DB Name') WITH PHYSICAL_ONLY
```

## It Has Found A Problem - What Do I Do?

You do have backups don't you? You might be lucky; a non-clustered index can be dropped and rebuilt but actual data, such as a clustered index, cannot.

By far the best option is to restore from a backup, but let's look at how you investigate which pages are affected and what type of data is affected:

Look at the output from DBCC CHECKDB. You may see something like this:

```
Msg 8928, Level 16, State 1, Line 1
Object ID 2088535921, index ID 0, partition ID
72345201021503994, alloc unit ID 72345201051571606 (type In-
row data): Page (1:94299) could not be processed. See other
errors for details.
Msg 8939, Level 16, State 98, Line 1
Table error: Object ID 2088535921, index ID 0, partition ID
72345201021503994, alloc unit ID 72345201051571606 (type In-
row data), page (1:94299). Test (IS_OFF (BUF_IOERR, pBUF-
>bstat)) failed.
CHECKDB found 0 allocation errors and 2 consistency errors in
table 'yourtable' (object ID 2088535921).
CHECKDB found 0 allocation errors and 2 consistency errors in
database 'yourdb'.
repair_allow_data_loss is the minimum repair level for the
errors found by DBCC CHECKDB (yourdb).
```

From this you can see what page is corrupted (1:94299)

The first thing to do is check if it is data in a heap, in a clustered index, or in a non-clustered index. In the above text you can see it is index ID 0. You could also examine the page (1:94299 in database 'yourdb') as follows:

```
DBCC TRACEON (3604, -1)
GO
DBCC PAGE('yourdb', 1, 94299, 3)
GO
```

In the output you will see something like:

*Metadata: IndexId = n*

If n is greater than 1 it is a non-clustered index and can safely be dropped and recreated. If n is 0 or 1 you have data corruption and need to perform one of the options described below.

### Restoring from a backup

If the recovery model is FULL (or BULK\_LOGGED, with some limitations), you can backup the tail of the log, perform a restore (with norecovery) from the last clean full backup, followed by subsequent log backups and finally the tail of the log.

If only a few pages are affected you have the option of selectively restoring only the bad pages, as follows:

```
RESTORE DATABASE yourdb PAGE = '1:94299'
FROM DISK = 'C:\yourdb.bak'
WITH NORECOVERY
```

If the recovery model is simple you don't have that option, and have to accept that a restore from the last full backup will result in subsequent transactions

being lost. In this case, or if you have no backups at all, you may decide that an automatic repair is the only option.

### Automatic Repair Options

First let me emphasise the importance of running a backup BEFORE you go any further.

Have a look at the output of the original CHECKDB. It will specify the minimum repair level.

#### *REPAIR\_REBUILD*

If the minimum repair level is REPAIR\_REBUILD you have been lucky. The syntax is

```
DBCC CHECKDB('DB Name', REPAIR_REBUILD)
```

#### *REPAIR\_ALLOW\_DATA\_LOSS*

This attempts to repair all errors. Sometimes the only way to repair an error is to deallocate the affected page and modify page links so that it looks like the page never existed. This has the desired effect of restoring the database's structural integrity but means that something has been deleted (hence the ALLOW\_DATA\_LOSS). There are likely to be issues with referential integrity, not to mention the important data that may now be missing.

The syntax is

```
DBCC CHECKDB('DB Name', REPAIR_ALLOW_DATA_LOSS)
```

Make sure you run DBCC CHECKCONSTRAINTS afterwards so you are aware of referential integrity issues and can take the appropriate action.

## And Finally

My original reason for writing this was to stress that the correct action when faced with corruption is nearly always to restore from a backup. Only use the automatic repair options as a last resort, and with full understanding of the damage this may do.

Just as important is that regular backups are an essential part of a DBA's responsibilities, and you should use the FULL recovery model with regular log backups for all but the most trivial databases.

DBCC CHECKDB is a powerful tool, but also very dangerous in the wrong hands.

Maybe instead of adding the REPAIR\_ALLOW\_DATA\_LOSS option, Microsoft should have created a separate DBCC command called:

```
DBCC  
DELETE_DATA_BECAUSE_I_COULDNT_BE_BOTHERED_TO_TAKE_A_BACKUP
```

A bit harsh perhaps, but at least people would know for sure what might happen when they hit F5!